

Assignment 2
CS289: Algorithmic Machine Learning, Fall 2016
Due: October 26, 10PM

Guidelines for submitting the solutions:

- The assignments need to be submitted on Gradescope. Make sure you follow all the instructions - they are simple enough that exceptions will not be accepted.
- To save my eyes (and perhaps, more importantly, a few of your points) please use a good scanner and/or digitize the scan as per the instructions on the class web page.
- Start each problem or sub-problem on a separate page even if it means having a lot of white-space and write/type in large font.
- The solutions need to be submitted by 10 PM on the due date. No late submissions will be accepted.
- Please adhere to the code of conduct outlined on the class page.

1. Exercise 3.4 from [BHK]. [2 points]
2. The goal of this exercise is to partially show one of the remarkable properties of SVD that we discussed in class. Let A be a $n \times d$ matrix and let $A = U\Sigma V^T$ be its singular-value decomposition where Σ is a $r \times r$ diagonal matrix with $\Sigma_{11} \geq \Sigma_{22} \geq \dots \geq \Sigma_{rr}$. Let $\sigma_i = \Sigma_{ii}$. Show that $\max_{v: \|v\|=1} \|Av\| \leq \sigma_1$. [2 points]
(Hint: Choose an appropriate orthonormal basis for \mathbb{R}^d and use the representation of vectors v in this basis to bound $\|Av\|$.)
3. Start by taking a picture of something around you with resolution about 1024x1024. You can view the image as 3 matrices (one for each color Red, Green, Blue). Now, compute the best rank- k approximation for these matrices and render the resulting image. Do this for $k = \{10, 20, 50, 100, 200\}$ (as I showed in class for a specific image). Your submission should be a) A screenshot of your code and b) The original image followed by the approximations (in order of increasing values of k). You can use MATLAB, R, Numpy/Scipy, or Julia for this. [2 points]
4. Let n be a large number (say $n = 10,000$) and k a small number (say 20). Let A be a sparse $n \times n$ matrix with density p (say 0.01). MATLAB (and other numerical analysis software) can exploit the sparsity of A to compute the top singular vector quite efficiently. Next, let U be a $n \times k$ matrix U and let $B = UU^T$. Computing the first right singular vector of B

directly on MATLAB would be quite slow. But, the first right singular vector of B is the same as the first left singular vector of U which you can compute much more efficiently.

Now, suppose we had to compute the first singular vector of $C = A + UU^T$ (such situations arise often in practice). Doing so directly would be quite expensive.

- (a) Give an algorithm based on the power iteration to approximate the top singular vector of C . Show that you can do this without even computing C explicitly and prove a bound on the per-iteration running-time (in terms of n, p, k) of the power iteration. [1 point]
- (b) You can use MATLAB, R, Numpy/Scipy, or Julia for this exercise. Generate a random $n \times n$ sparse matrix with density p for $n = 10,000$, $p = 0.01$. (In MATLAB, you can do this by `sprand(n,n,p)`.) Generate a random $n \times k$ matrix U . Compute the top singular vector/value of $C = A + UU^T$ using the in-built MATLAB command and time the operation. Implement the power iteration to compute the top singular vector of $C = A + UU^T$ (say for 10, 20, 30, ..., 100 iterations) and time the algorithm. Is there a difference in speed?

Your submission should be the following items. [3 points]

- i. A screenshot of your code for the power iteration (only the important part - this should only take a few lines).
- ii. A table or figure showing the run-times of the in-built MATLAB command and the run-times for different number of iterations obtained by averaging over 10 runs (as A, U are random - it is better to consider multiple runs).
- iii. A figure plotting the number of iterations (on the x-axis) with the error on the y-axis averaged over 10 runs.